

Exercise 1: GD on a 1D Function
Exercise 2: Backtracking Line Search
Exercise 3: GD in 2D
Exercise 4: Exact Line Search vs Backtracking
Exercise 5: Gradient Descent on the Rosenbrock Function

Homework 1: Gradient Descent

Exercise 1: GD on a 1D Function

Consider the 1-dimensional function

$$\mathcal{L}(\Theta) = (\Theta - 3)^2 + 1.$$

1. Compute the Gradient of $\mathcal{L}(\Theta)$ explicitly.
2. Implement the Gradient Descent to optimize $\mathcal{L}(\Theta)$ following what we introduced on the theoretical sections.
3. Test three different constant step sizes:
 - $\eta = 0.05$
 - $\eta = 0.2$
 - $\eta = 1.0$
4. For each choice:
 - Plot the sequence $\Theta^{(k)}$ on the real line (i.e. draw a line and represent, on top of it, the position of all the successive elements $\Theta^{(k)}$).
 - Plot the function values $\mathcal{L}(\Theta^{(k)})$ vs iteration.
 - Comment on convergence, oscillations, and divergence.
5. Relate your observations to the discussion in class about:
 - step-size being too small / too large,
 - the role of convexity,
 - how the "just right" step size leads to fast convergence.

Hint: This function is strictly convex with a unique minimizer at $\Theta^* = 3$.

Exercise 2: Backtracking Line Search

Consider the non-convex function

$$\mathcal{L}(\Theta) = \Theta^4 - 3\Theta^2 + 2.$$

1. Implement Gradient Descent with Backtracking, using the Armijo condition, considering the `backtracking(...)` function from class.
2. Test different initial points:
 - $\Theta_0 = -2$
 - $\Theta_0 = 0.5$
 - $\Theta_0 = 2$
3. For each starting point, plot:
 - the function curve $\mathcal{L}(\Theta)$ in 1D in the domain $[-3, 3]$,
 - the trajectory of the iterates $\Theta^{(k)}$ overlaid on the curve.
4. Discuss:
 - Why different initializations converge to different minima.
 - How backtracking automatically chooses a suitable step size at each iteration.
 - Situations where constant step size would fail.

Exercise 3: GD in 2D

Consider the quadratic function:

$$\mathcal{L}(\Theta) = \frac{1}{2} \Theta^T A \Theta, \quad A = \begin{bmatrix} 1 & 0 \\ 0 & 25 \end{bmatrix}.$$

1. Implement Gradient Descent in 2D, as seen in class.
2. Plot the level sets of L (using the helper code provided in the lecture notes) and overlay the iterates $\Theta^{(k)}$ for:
 - $\eta = 0.02$,
 - $\eta = 0.05$ (borderline),
 - $\eta = 0.1$ (diverging). To overlay the iterates on the plot, simply add a `plt.plot(...)` function taking as input the coordinates $\Theta_1^{(k)}$ and $\Theta_2^{(k)}$ for each k .
3. Comment on:
 - the elongated ellipses produced by ill-conditioning,
 - the gradient direction compared to the level set lines,
 - zig-zag behaviour for large condition numbers,
 - the relation to the lecture discussion on slow convergence in narrow valleys.

Exercise 4: Exact Line Search vs Backtracking

Let

$$\mathcal{L}(\Theta) = \frac{1}{2} \Theta^T A \Theta, \quad A = \begin{bmatrix} 5 & 0 \\ 0 & 2 \end{bmatrix}.$$

1. For GD with exact line search i.e. selecting the optimal step, you get:

$$\eta_k^* = \frac{\bar{g}_k^T g_k}{\bar{g}_k^T A g_k}, \quad g_k = A \Theta^{(k)}.$$
2. For GD with backtracking, use the implementation provided in your notes.
3. Starting from $\Theta_0 = (3, 3)^T$:
 - Plot trajectories on the level-set map.
 - Plot the loss $\mathcal{L}(\Theta^{(k)})$ vs iteration for both methods.
4. Compare:
 - speed of convergence,
 - smoothness of step-sizes.

Exercise 5: Gradient Descent on the Rosenbrock Function

The 2-dimensional Rosenbrock function is defined as:

$$f(\Theta) = (1 - \Theta_1)^2 + 100(\Theta_2 - \Theta_1^2)^2,$$

with a unique global minimum at:

$$\Theta^* = (1, 1), \quad \mathcal{L}(\Theta_1^*, \Theta_2^*) = 0.$$

Despite having only one minimum, its landscape is **extremely challenging**:

- The valley is long and curved.
- The condition number varies drastically across the region.
- Gradients can point almost orthogonally to the valley direction.
- Gradient Descent may zig-zag and make very slow progress toward the solution.

This makes Rosenbrock a perfect testbed to study the difficulties of pure Gradient Descent.

1. Implement the function and its gradient to run the Gradient Descent algorithm. Test its performance using both **constant step size** and the **backtracking algorithm**.

2. Plot the level sets of \mathcal{L} : use the standard 2D contour plotting approach practiced earlier.

Ensure that you plot:

- A wide range (e.g. $\Theta_1 \in [-2, 2]$, $\Theta_2 \in [-1, 3]$)
- A reasonable number of contour lines to visualize the narrow valley.

1. Choose multiple initial points, testing at least the following four:

- $\Theta^{(0)} = (-1.5, 2)$
- $\Theta^{(0)} = (-1, 0)$
- $\Theta^{(0)} = (0, 2)$
- $\Theta^{(0)} = (1.5, 1.5)$

Run both:

- GD with constant step size (try $\eta = 10^{-3}, 10^{-4}, 10^{-5}$),
- GD with backtracking line search.

2. Visualize and analyze the trajectories: For each initialization and each training method:

1. Plot the sequence of iterates $(\Theta_1^{(k)}, \Theta_2^{(k)})$ as a path on the level-set map.

2. Plot the loss curve $L(\Theta^{(k)})$ vs iteration.

3. Observe and explain:

- whether the method enters the valley,
- how long it takes to “turn” correctly along the valley direction,
- whether the method zig-zags,
- whether step sizes become too small (with constant η) or too large (divergence).

< Previous
[Unsupervised Learning via Linear Algebra](#)

Next >
[Homework 2: Stochastic Gradient Descent](#)